Course 2: Tokenization

What is tokenization?

Turning text...

```
I love playing soccer!
```

...into *tokens*

```
['I', 'love', 'play', 'ing', 'soccer', '!']
```

Historical Notions

Tokenization Origins

The word token comes from linguistics

" *non-empty contiguous sequence of graphemes or phonemes in a document* "

$\approx$

Split text on blanks

Tokenization Origins

```
old_tokenize("I love playing soccer!") = ['I', 'love', 'playing', 'soccer!']
```

- Different from *word-forms* ⚠️
  - ○ *damélo → da/mé/lo (=give/me/it)*

Tokenization Origins

Natural language is split into…

- Sentences, utterances, documents… (*macroscopical*)
  that are split into…
  - Tokens, word-forms… (*microscopical*)

→ Used for linguistic tasks (POS tagging, syntax parsing,…)

Tokenization & ML

Machine Learning relies on **sub-word** tokenization:

- Gives better performance
- **Fixed-size vocabulary** often required
    - Out-Of-Vocabulary (OOV) issue

Tokenization & ML

Evolution of modeling complexity w.r.t. the sequence length n

Model Type	Year	Complexity
Tf-Idf	1972	$O(1)$
RNNs	~1985	$O(n)$
Transformers	2017	$O(n^2)$

→ Long sequences (e.g. character-level) are prohibitive

Modern framework

- **Pre-tokenization**

  `"I'm lovin' it!" -> ["i", "am", "loving", "it", "!"]`

  - Normalization
    - Rules around punctuation (`_:_` , `_!` , …)
    - Spelling correction (`"imo" -> "in my opinion"`)
    - Named entities (`"covid" -> "COVID-19"`)
    - …
  - Rule-based segmentation
    - Blanks, punctuation, …

Modern framework

- **Tokenization** `-> ["i", "am", "lov", "##ing", "it", "!"]`
  - Split units at subword level
  - Fixed vocabulary
  - **Trained** on text samples
  - Used in inference mode at *pre-processing* time

Sub-word Tokenization

Granularity

Vocabulary size (vertical axis)

Sequence lengths (horizontal axis)

extraordinarily

extra | ordinarily

extra | ordinar | ily

ex | tra | ord | inar | ily

e | x | t | r | a …

Granularity

→ Trade-off between short sequences and reasonable vocabulary size

Fertility
For a string sequence S:

$$\text{fertility}(S) = \frac{\#\text{ tokens}}{\#\text{ words}}$$

Algorithms

Byte-Pair Encoding (BPE)

Let's encode "*aaabdaaabac*" in an optimized way:

- Observed pairs: {*aa*, *ab*, *bd*, *da*, *ba*, *ac*}
- Observed **occurences**: {***aa*: 4**, *ab*: 2, *bd*: 1, *da*: 1, *ba*: 1, *ac*: 1}
- Set *X = aa*
- Encode *aaabdaaabac* → *XabdXabac*
- Start again from *XabdXabac*

Byte-Pair Encoding (BPE)

(current rules: *aa* → *X*)

Let's encode "*XabdXabac*" in an optimized way:

- Observed pairs: {*Xa*, *ab*, *bd*, *dX*, *ba*, *ac*}
- Observed **occurences**: {**Xa: 2**, **ab: 2**, *bd*: 1, *dX*: 1, *ba*: 1, *ac*: 1}
- Set *Y = ab*
- Encode *XabdXabac* → *XYdXYac*
- Start again from *XYdXYac*

Byte-Pair Encoding (BPE)

(current rules: *aa* → *X*, *ab* → *Y*)

Let's encode "*XYdXYac*" in an optimized way:

- Observed pairs: {*XY*, *Yd*, *dX*, *Ya*, *ac*}
- Observed occurences: {**XY: 2**, *Yd*: 1, *dX*: 1, *Ya*: 1, *ac*: 1}
- Set *Z = XY*
- Encode *XYdXYac* → *ZdZac*
- Start again from *ZdZac*

Byte-Pair Encoding (BPE)

(current rules: $aa \rightarrow X$, $ab \rightarrow Y$, $XY \rightarrow Z$)

Let's encode "*ZdZac*" in an optimized way:

- Observed pairs: {*Zd, dZ, Za, ac*}
- Observed occurences: {*Zd*: 1, *dZ*: 1, *Za*: 1, *ac*: 1}
- **All pairs are unique => END**

Byte-Pair Encoding (BPE)

Final encoding: $aaabdaaabac \rightarrow ZdZac$

with **merge rules**:

1. $aa \rightarrow X$
2. $ab \rightarrow Y$
3. $XY \rightarrow Z$

Decoding: follow merge rules in opposite order

BPE Training - pre-tokenization

```
training_sentences = [
    "Education is very important!",
    "A cat and a dog live on an island",
    "We'll be landing in Cabo Verde",
]
```

=>

```
pretokenized = ["education_", "is_", "very_", "important_", "!_", "a_",
        "cat_", "and_", "a_", "dog_", "live_", "on_", "an_", "island_",
        "we", "'", "ll_", "be_", "landing_", "in_", "cabo_" "Verde_"
]
```

BPE Training - iteration 1

```
tokenized = [
    ['e', 'd', 'u', 'c', 'a', 't', 'i', 'o', 'n', '_'], ..., ['i', 'm', 'p', 'o', 'r', 't', 'a', 'n', 't', '_'], ['!', '_'],
    ['a', '_'], ['c', 'a', 't', '_'], ['a', 'n', 'd', '_'],  ..., ['o', 'n', '_'], ['a', 'n', '_'], ['i', 's', 'l', 'a', 'n', 'd', '_'],
    ['w', 'e'], ['"'], ['l', 'l', '_'], ['b', 'e', '_'], ['l', 'a', 'n', 'd', 'i', 'n', 'g', '_'], ..., ['v', 'e', 'r', 'd', 'e', '_']
]
```

→ Most common pair: "an"

```
tokenized = [
    ['e', 'd', 'u', 'c', 'a', 't', 'i', 'o', 'n', '_'], ..., ['i', 'm', 'p', 'o', 'r', 't', 'an', 't', '_'], ['!', '_'],
    ['a', '_'], ['c', 'a', 't', '_'], ['an', 'd', '_'],  ..., ['o', 'n', '_'], ['an', '_'], ['i', 's', 'l', 'an', 'd', '_'],
    ['w', 'e'], ['"'], ['l', 'l', '_'], ['b', 'e', '_'], ['l', 'an', 'd', 'i', 'n', 'g', '_'], ..., ['v', 'e', 'r', 'd', 'e', '_']
]
```

BPE Training - iteration 2

```
tokenized = [
    ['e', 'd', 'u', 'c', 'a', 't', 'i', 'o', 'n', '_'], ..., ['i', 'm', 'p', 'o', 'r', 't', 'an', 't', '_'], ['!', '_'],
    ['a', '_'], ['c', 'a', 't', '_'], ['an', 'd', '_'],  ..., ['o', 'n', '_'], ['an', '_'], ['i', 's', 'l', 'an', 'd', '_'],
    ['w', 'e'], ['"'], ['l', 'l', '_'], ['b', 'e', '_'], ['l', 'an', 'd', 'i', 'n', 'g', '_'], ..., ['v', 'e', 'r', 'd', 'e', '_']
]
```

→ Most common pair: `"ca"`

```
tokenized = [
    ['e', 'd', 'u', 'ca', 't', 'i', 'o', 'n', '_'], ..., ['i', 'm', 'p', 'o', 'r', 't', 'an', 't', '_'], ['!', '_'],
    ['a', '_'], ['ca', 't', '_'], ['an', 'd', '_'],  ..., ['o', 'n', '_'], ['an', '_'], ['i', 's', 'l', 'an', 'd', '_'],
    ['w', 'e'], ['"'], ['l', 'l', '_'], ['b', 'e', '_'], ['l', 'an', 'd', 'i', 'n', 'g', '_'], ..., ['v', 'e', 'r', 'd', 'e', '_']
]
```

BPE Training - iteration 14 (final)

```
tokenized = [
    ['e', 'd', 'u', 'cat', 'i', 'on_'], ['is', '_'], ['ver', 'y', '_'], ['i', 'm', 'p', 'o', 'r', 't', 'an', 't', '_'], ['!', '_'],
    ['a_'], ['cat', '_'], ['and_'], ['a_'], ..., ['on_'], ['an', '_'], ['is', 'l', 'and_'],
    ['w', 'e'], ["'"], ..., ['l', 'and', 'i', 'n', 'g_'], ['i', 'n_'], ['ca', 'b', 'o', '_'], ['ver', 'd', 'e_']
]
```
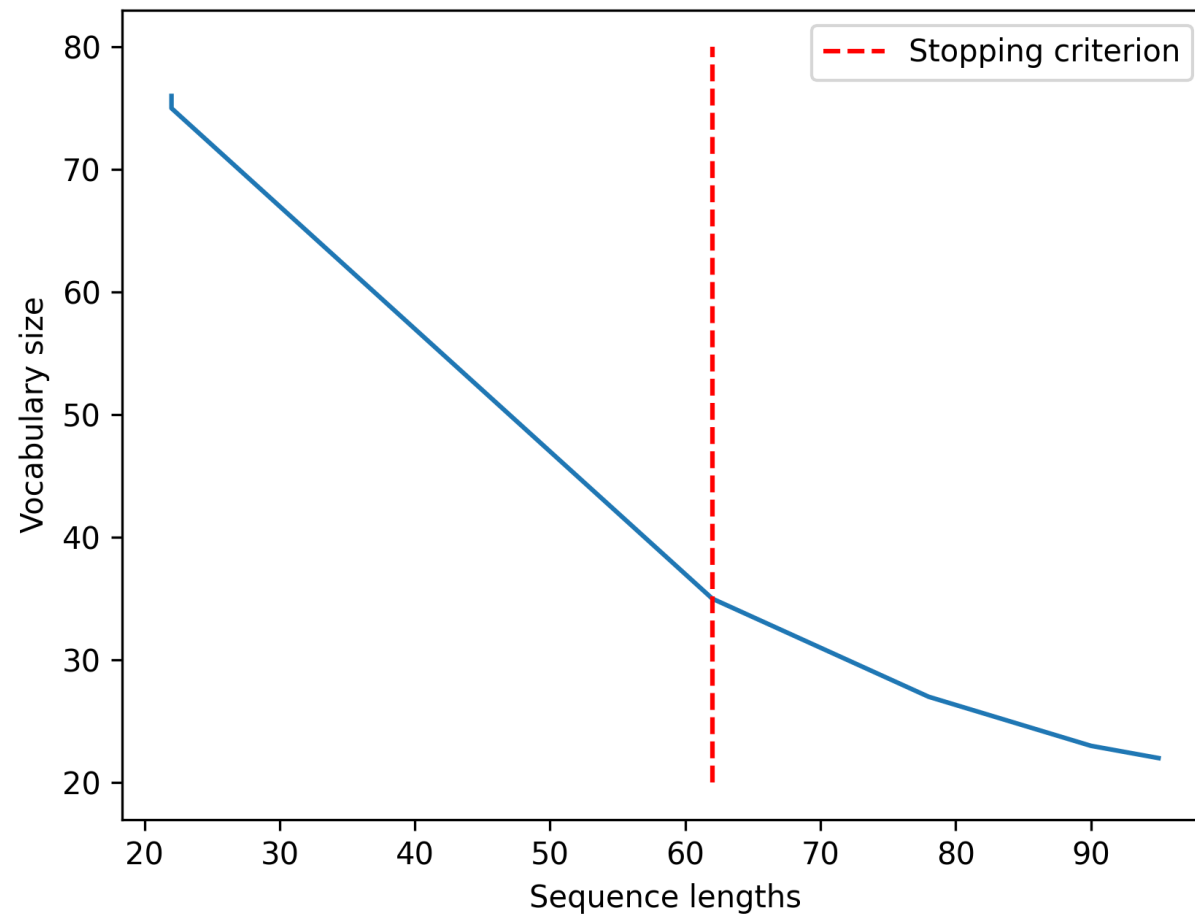
"Created" tokens:

```
['an', 'ca', 'n_', 've', 'and', 'cat', 'on_', 'is', 'ver', 'a_', 'and_', 'g_', 'e_']
```

→ English common words (a, and, on, is, ...)

→ `and` vs `and_`

BPE Training - iteration 14 (final)

```
tokenized = [
    ['e', 'd', 'u', 'cat', 'i', 'on_'], ['is', '_'], ['ver', 'y', '_'], ['i', 'm', 'p', 'o', 'r', 't', 'an', 't', '_'], ['!', '_'],
    ['a_'], ['cat', '_'], ['and_'], ['a_'], ..., ['on_'], ['an', '_'], ['is', 'l', 'and_'],
    ['w', 'e'], ["'"], ..., ['l', 'and', 'i', 'n', 'g_'], ['i', 'n_'], ['ca', 'b', 'o', '_'], ['ver', 'd', 'e_']
]
```

"Created" tokens:

```
['an', 'ca', 'n_', 've', 'and', 'cat', 'on_', 'is', 'ver', 'a_', 'and_', 'g_', 'e_']
```

→ English common words (a, and, on, is, ...)
→ and vs and_

BPE - Granularity

WordPiece

- Based on merge rules too
- Initial processing is different:

BPE:

```
["education", "is"] => [["e", "d", "u", ..., "n", "_"], ["i", "s", "_"]]
```

WordPiece:

```
["education", "is"] => [["e", "##d", "##u", "##c",...], ["i", "##s"]]
```

WordPiece

- Pairs are scored using this score function:

$$S((t_1, t_2)) = \frac{freq(t_1 t_2)}{freq(t_1)\,freq(t2)}$$

- if t_1 and t_2 are common, less likely to merge
  - ex: *dream/##ing* → not merged
- if t_1 and t_2 are rare but $t_1 t_2$ is common, more likely to merge
  - ex: *pulv/##erise* → *pulverise*

Unigram

Unigram is working in the opposite direction:

- Start from a (too) big subword vocabulary

- Gradually eliminate tokens that won't be missed 👋

- **Score** all possible segmentations and take max:
  - Ex: *brew*
    - S(*b*/*r*/*e*/*w*) → P(*b*) x P(*r*) x P(*e*) x P(*w*) = 0.024
    - S(*br*/*e*/*w*) → P(*br*) x P(*e*) x P(*w*) = 0.031
    - ...

Unigram - Inference

⚠️ A string of length n has $O(2^n)$ possible segmentations ⚠️

$\rightarrow$ Unigram is using the Viterbi algorithm:

- Observation: $P(S^*(w_{:i+1})) \leq P(S^*(w_{:i}))$ (S^*: best segmentation)

Unigram - Inference

Example: *email*

- Starting from letter *e*
  - For all <u>ending letters</u>, what is the best segmentation if last token starts from *e*?
    - S(*e*) = 0.15
    - S(*em*) = 0.02
    - ...
    - S(*email*) = 0.001

Unigram - Inference

Example: *email*

- Starting from letter *m*
  - For all <u>ending letters</u>, what is the best segmentation if last token starts from *m*?
    - S(e/m) = 0.1

      ...

    - S(e/$mail$) = 0.2
- Remark: we've seen S(em) and S(e/m) $\rightarrow$ we know the best segmentation that ends at *m* !

Unigram - Inference

Example: *email*

- Starting from letter *a*
  - For all <u>ending letters</u>, what is the best segmentation if last token starts from *a*? **(hence after *e* / *m*)**
    - S(*e* / *m* / *a*) = 0.023

      ...

    - S(*e* / *m*/ *ail*) = 0.001
- Remark: we've seen S(*ema*), ..., S(*e* / *m*/ *a*) → we know the best segmentation that ends at *a* ! (here: *e* / *ma* is best)

Unigram - Inference

Example: *email*

- Starting from letter *i*
  - For all <u>ending letters</u>, what is the best segmentation if last token starts from *i* ? (hence after *e* / *ma*)
    - S(*e* / *ma* / *i*) = 0.004
    - S(*e* / *ma* / *il*) = 0.03
- Remark: we only have 2 candidates left! (here: *ema* / *i* is best)

Unigram - Inference

Example: *email*

- Starting from letter *l*
  - For all <u>ending letters</u>, what is the best segmentation if last token starts from *l*? (hence after *ema/i*)
    - S(*ema/i/l*) = 0.002

Takeaway: At each *start* position, we know what the best segmentation up to *start* is => we just need to explore after *start*

Unigram - Training

- Start from a very big vocabulary

- Infer on all pre-tokenized units $w \in W$ and get total score as:

Unigram - Training

- Start from a very big vocabulary

- Infer on all pre-tokenized units $w \in W$ and get total score as:

$$score(V, W) = \sum_{w=(t_1...t_n) \in W} -\log(P(t_1) \times \ldots \times P(t_n))$$

- For all token t, compute $score(V - \{t\}, W)$

- Remove the token that **least decreases** the score when removed

- Iterate 🔁 (until you have desired vocabulary size)

Limits & Alternatives

Limits

- Fixed vocabulary...
    - ... leads to OOV (out-of-vocabulary)
    - ... scales poorly to 100+ languages (and scripts)
    - ... can cause over-segmentation
    - ... is not robust to misspellings

```
bpe("artificial intelligence is real bro") => 'artificial', 'intelligence', 'is', 'real', 'bro'
```

```
bpe("aritificial inteligense is reaal bro") =>
'ari', '##ti', '##fi', '##cial', 'intel', '##igen', '##se', 'is', 're', '##aa', '##l', 'bro'
```

Alternatives - BPE dropout

→ Randomly removes part of the vocabulary during training


```
u-n-r-e-l-a-t-e-d
u-n re-l-a-t-e-d
u-n re-l-at-e-d
u-n re-l-at-ed
un re-l-at-ed
un re-l-ated
un rel-ated
un-related
unrelated
        (a)
```

```
u-n_r-e-l-a_t-e_d        u-n-r-e-l-a_t-e-d        u-n_r_e_l-a-t-e-d
u-n re-l_a-t-e_d         u_n re_l-a-t-e-d         u-n-r_e-l-at-e-d
u-n re_l-at-e_d          u_n re-l-at-e-d          u-n-r_e-l_at_ed
un re-l-at-e-d           u_n re-l-ate_d           un-r-e-l-at-ed
un re_l-at-ed            u_n rel-ate-d            un re-l_at-ed
un re-lat-ed             u_n relate_d             un re-l-ated
un relat_ed                                       un rel_ated
        (b)
```

Figure 1: Segmentation process of the word '*unrelated*' using (a) BPE, (b) *BPE-dropout*. Hyphens indicate possible merges (merges which are present in the merge table); merges performed at each iteration are shown in green, dropped – in red.
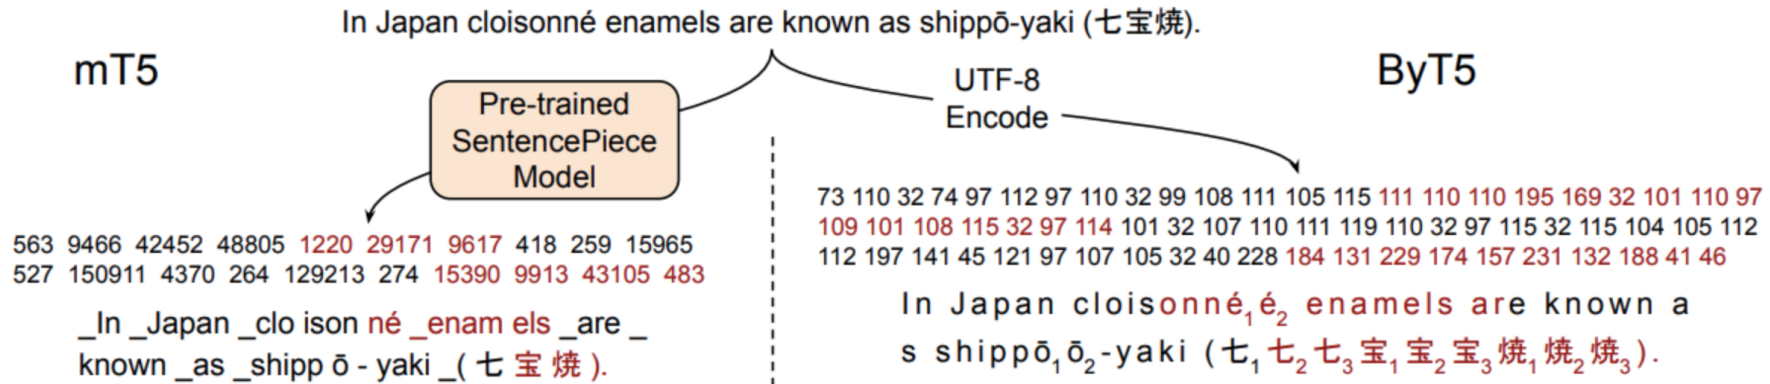
=> makes models more robust to misspellings

Alternatives - CharacterBERT
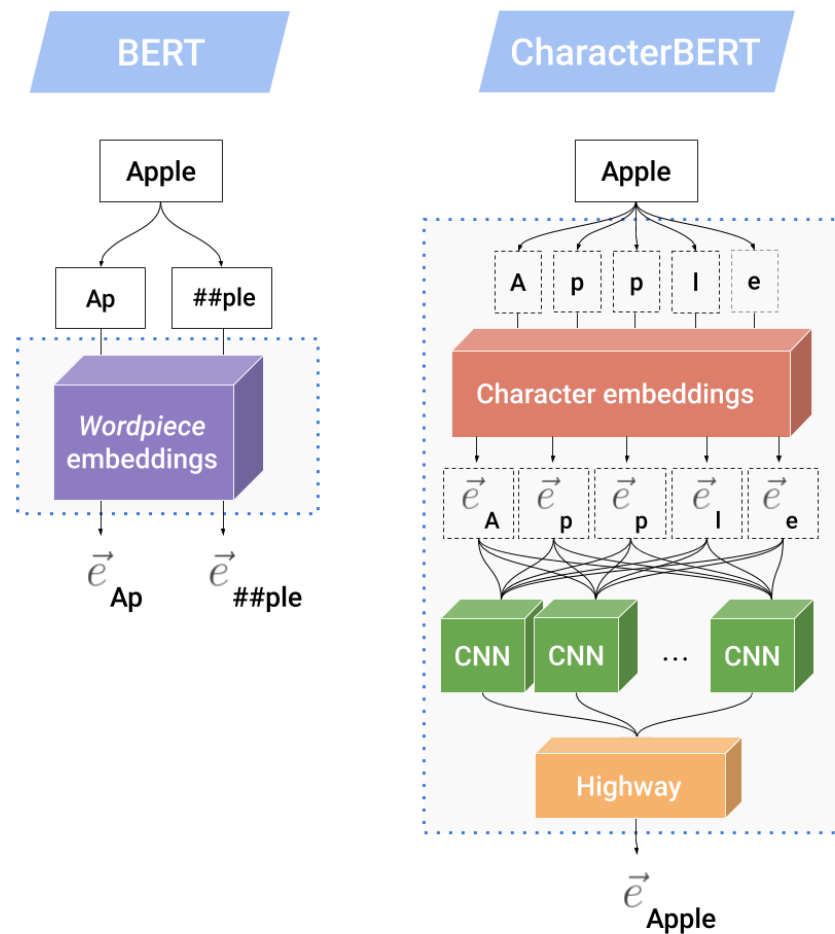
Alternatives - ByT5

- Gives directly bytes (~characters) as inputs to the model

=> more robust and data efficient BUT ~10 times slower and more hardware consumption

Neural tokenization - CANINE

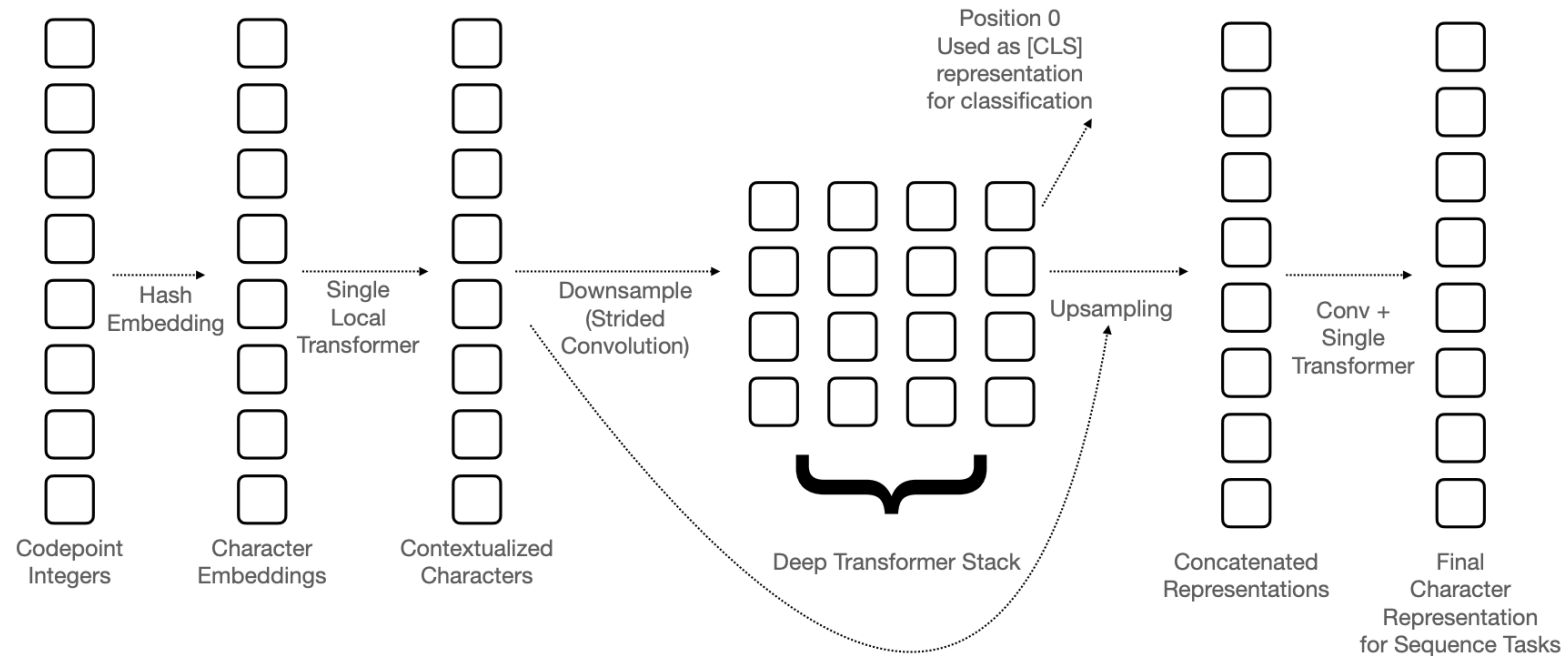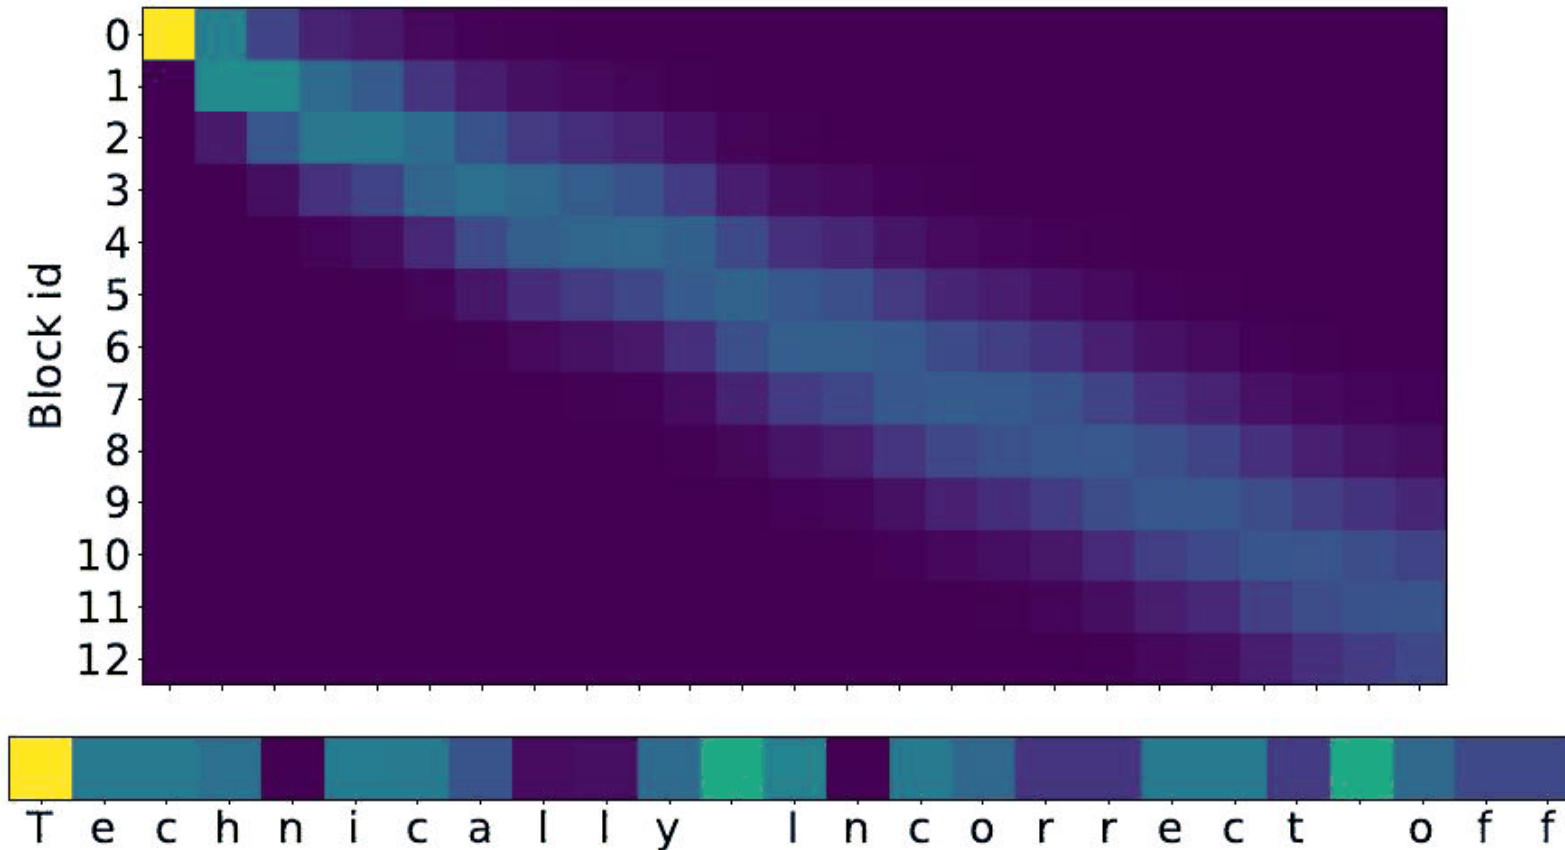- Downsamples characters into $4\times$ smaller sequences

Figure 1: CANINE neural architecture.

Neural tokenization - MANTa

- Allows the language model to learn its *own* mapping

Takeaways

- Tokenization: Art of splitting sentences/words into meaningful smaller units
- In ML: subword tokenization is (very) commonly used
- Three main algorithms
  - **BPE**: iteratively learn most frequent merges
  - **WordPiece**: BPE with adjusted frequency score
  - **Unigram**: Start big and remove tokens that won't be missed
- When facing noisy and/or complex text, alternatives exist